

Code

```
#include <iostream>
#include <vector>
#include <queue>
#include <ctime>
#include <omp.h>

using namespace std;

// Function to perform BFS from a given vertex
void bfs(int startVertex, vector<bool> &visited, vector<vector<int>> &graph)
{
    // Create a queue for BFS
    queue<int> q;

    // Mark the start vertex as visited and enqueue it
    visited[startVertex] = true;
    q.push(startVertex);

    // Loop until the queue is empty
    while (!q.empty())
    {
        // Dequeue a vertex from the queue
        int v = q.front();
        q.pop();

        // Enqueue all adjacent vertices that are not visited
        #pragma omp parallel for
        for (int i = 0; i < graph[v].size(); i++)
        {
            int u = graph[v][i];
            #pragma omp critical
            {
                if (!visited[u])
                {
                    visited[u] = true;
                    q.push(u);
                }
            }
        }
    }
}

// Parallel Breadth-First Search
void parallelBFS(vector<vector<int>> &graph, int numCores)
{
    int numVertices = graph.size();
    vector<bool> visited(numVertices, false); // Keep track of visited vertices

    double startTime = omp_get_wtime(); // Start timer

    // Perform BFS from all unvisited vertices using specified number of cores
```

```

#pragma omp parallel for num_threads(numCores)
for (int v = 0; v < numVertices; v++)
{
    if (!visited[v])
    {
        bfs(v, visited, graph);
    }
}

double endTime = omp_get_wtime(); // End timer

cout << "Number of cores used: " << numCores << endl;
cout << "Time taken: " << endTime - startTime << " seconds" << endl;
cout << "-----" << endl;
}

int main()
{
    // Generate a random graph with 10,000 vertices and 50,000 edges
    int numVertices = 10000;
    int numEdges = 50000;
    vector<vector<int>> graph(numVertices);
    srand(time(0));
    for (int i = 0; i < numEdges; i++)
    {
        int u = rand() % numVertices;
        int v = rand() % numVertices;
        graph[u].push_back(v);
        graph[v].push_back(u);
    }

    // Array containing number of cores
    int numCoresArr[] = {1, 2, 3, 4, 5, 6, 7, 8};

    // Loop over different number of cores and execute parallel BFS
    for (int i = 0; i < sizeof(numCoresArr) / sizeof(numCoresArr[0]); i++)
    {
        int numCores = numCoresArr[i];
        cout << "Running parallel BFS with " << numCores << " core(s)..." << endl;
        parallelBFS(graph, numCores);
    }

    return 0;
}

```

Output

```
Windows PowerShell
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> g++ dfs.cpp -o dfs -fopenmp -pthread
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> ./dfs
Running parallel DFS with 1 core(s)...
Number of cores used: 1
Time taken: 10.921 seconds
-----
Running parallel DFS with 2 core(s)...
Number of cores used: 2
Time taken: 0.0769999 seconds
-----
Running parallel DFS with 3 core(s)...
Number of cores used: 3
Time taken: 0.0639999 seconds
-----
Running parallel DFS with 4 core(s)...
Number of cores used: 4
Time taken: 0.0469999 seconds
-----
Running parallel DFS with 5 core(s)...
Number of cores used: 5
Time taken: 0.043 seconds
-----
Running parallel DFS with 6 core(s)...
Number of cores used: 6
Time taken: 0.0400002 seconds
-----
Running parallel DFS with 7 core(s)...
Number of cores used: 7
Time taken: 0.0380001 seconds
-----
Running parallel DFS with 8 core(s)...
Number of cores used: 8
Time taken: 0.0370002 seconds
-----
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> |
```

Code

```
#include <iostream>
#include <vector>
#include <stack>
#include <ctime>
#include <omp.h>

using namespace std;

// Function to perform DFS from a given vertex
void dfs(int startVertex, vector<bool> &visited, vector<vector<int>> &graph)
{
    // Create a stack for DFS
    stack<int> s;

    // Mark the start vertex as visited and push it onto the stack
    visited[startVertex] = true;
    s.push(startVertex);

    // Loop until the stack is empty
    while (!s.empty())
    {
        // Pop a vertex from the stack
        int v = s.top();
        s.pop();

        // Push all adjacent vertices that are not visited onto the stack
        #pragma omp parallel for
        for (int i = 0; i < graph[v].size(); i++)
        {
            int u = graph[v][i];
            #pragma omp critical
            {
                if (!visited[u])
                {
                    visited[u] = true;
                    s.push(u);
                }
            }
        }
    }
}

// Parallel Depth-First Search
void parallelDFS(vector<vector<int>> &graph, int numCores)
{
    int numVertices = graph.size();
    vector<bool> visited(numVertices, false); // Keep track of visited vertices

    double startTime = omp_get_wtime(); // Start timer

    // Perform DFS from all unvisited vertices using specified number of cores
```

```

#pragma omp parallel for num_threads(numCores)
for (int v = 0; v < numVertices; v++)
{
    if (!visited[v])
    {
        dfs(v, visited, graph);
    }
}

double endTime = omp_get_wtime(); // End timer

cout << "Number of cores used: " << numCores << endl;
cout << "Time taken: " << endTime - startTime << " seconds" << endl;
cout << "-----" << endl;
}

int main()
{
    // Generate a random graph with 10,000 vertices and 50,000 edges
    int numVertices = 10000;
    int numEdges = 50000;
    vector<vector<int>> graph(numVertices);
    srand(time(0));
    for (int i = 0; i < numEdges; i++)
    {
        int u = rand() % numVertices;
        int v = rand() % numVertices;
        graph[u].push_back(v);
        graph[v].push_back(u);
    }

    // Array containing number of cores
    int numCoresArr[] = {1, 2, 3, 4, 5, 6, 7, 8};

    // Loop over different number of cores and execute parallel DFS
    for (int i = 0; i < sizeof(numCoresArr) / sizeof(numCoresArr[0]); i++)
    {
        int numCores = numCoresArr[i];
        cout << "Running parallel DFS with " << numCores << " core(s)..." << endl;
        parallelDFS(graph, numCores);
    }

    return 0;
}

```

Output

```
Windows PowerShell
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> g++ bfs.cpp -o dfs -fopenmp -pthread
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> ./bfs
Running parallel BFS with 1 core(s)...
Number of cores used: 1
Time taken: 11.01 seconds
-----
Running parallel BFS with 2 core(s)...
Number of cores used: 2
Time taken: 0.076 seconds
-----
Running parallel BFS with 3 core(s)...
Number of cores used: 3
Time taken: 0.0679998 seconds
-----
Running parallel BFS with 4 core(s)...
Number of cores used: 4
Time taken: 0.0539999 seconds
-----
Running parallel BFS with 5 core(s)...
Number of cores used: 5
Time taken: 0.0469999 seconds
-----
Running parallel BFS with 6 core(s)...
Number of cores used: 6
Time taken: 0.043 seconds
-----
Running parallel BFS with 7 core(s)...
Number of cores used: 7
Time taken: 0.0419998 seconds
-----
Running parallel BFS with 8 core(s)...
Number of cores used: 8
Time taken: 0.0409999 seconds
-----
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> |
```

Code

```
#include <omp.h>
#include <stdlib.h>

#include <array>
#include <chrono>
#include <functional>
#include <iostream>
#include <string>
#include <vector>

using std::chrono::duration_cast;
using std::chrono::high_resolution_clock;
using std::chrono::milliseconds;
using namespace std;

void s_bubble(int *, int);
void p_bubble(int *, int);
void swap(int &, int &);

void s_bubble(int *a, int n)
{
    for (int i = 0; i < n; i++)
    {
        int first = i % 2;
        for (int j = first; j < n - 1; j += 2)
        {
            if (a[j] > a[j + 1])
            {
                swap(a[j], a[j + 1]);
            }
        }
    }
}

void p_bubble(int *a, int n)
{
    for (int i = 0; i < n; i++)
    {
        int first = i % 2;
#pragma omp parallel for shared(a, first) num_threads(16)
        for (int j = first; j < n - 1; j += 2)
        {
            if (a[j] > a[j + 1])
            {
                swap(a[j], a[j + 1]);
            }
        }
    }
}

void swap(int &a, int &b)
```

```

{
    int test;
    test = a;
    a = b;
    b = test;
}

int bench_traverse(std::function<void()> traverse_fn)
{
    auto start = high_resolution_clock::now();
    traverse_fn();
    auto stop = high_resolution_clock::now();

    // Subtract stop and start timepoints and cast it to required unit.
    // Predefined units are nanoseconds, microseconds, milliseconds, seconds,
    // minutes, hours. Use duration_cast() function.
    auto duration = duration_cast<milliseconds>(stop - start);

    // To get the value of duration use the count() member function on the
    // duration object
    return duration.count();
}

int main(int argc, const char **argv)
{
    if (argc < 2)
    {
        cout << "Specify array length.\n";
        return 1;
    }
    int *a, n;
    n = stoi(argv[1]);
    a = new int[n];

    for (int i = 0; i < n; i++)
    {
        a[i] = rand() % n;
    }

    int *b = new int[n];
    copy(a, a + n, b);
    cout << "Generated random array of length " << n << "\n\n";

    int sequentialTime = bench_traverse([&]
                                         { s_bubble(a, n); });

    omp_set_num_threads(16);
    int parallelTime = bench_traverse([&]
                                       { s_bubble(a, n); });

    float speedUp = (float)sequentialTime / parallelTime;
    float efficiency = speedUp / 16;

```



```
cout
    << "Sequential Bubble sort: " << sequentialTime << "ms\n";

cout << "Parallel (16) Bubble sort: " << parallelTime << "ms\n";

cout << "Speed Up: " << speedUp << "\n";

cout << "Efficiency: " << efficiency << "\n";

return 0;
}
```

Output

```
Windows PowerShell
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> g++ bubble.cpp -o bubble -fopenmp
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> ./bubble 100000
Generated random array of length 100000

Sequential Bubble sort: 36615ms
Parallel (16) Bubble sort: 15881ms
Speed Up: 2.30559
Efficiency: 0.144099
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> |
```

Code

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <iostream>

using namespace std;

#define ARRAY_SIZE 5000

void merge(int arr[], int left[], int left_size, int right[], int right_size)
{
    int i = 0, j = 0, k = 0;
    while (i < left_size && j < right_size)
    {
        if (left[i] <= right[j])
        {
            arr[k] = left[i];
            i++;
        }
        else
        {
            arr[k] = right[j];
            j++;
        }
        k++;
    }
    while (i < left_size)
    {
        arr[k] = left[i];
        i++;
        k++;
    }
    while (j < right_size)
    {
        arr[k] = right[j];
        j++;
        k++;
    }
}

void merge_sort(int arr[], int size)
{
    if (size < 2)
    {
        return;
    }
    int mid = size / 2;
    int left[mid], right[size - mid];
    for (int i = 0; i < mid; i++)
    {
```

```

        left[i] = arr[i];
    }
    for (int i = mid; i < size; i++)
    {
        right[i - mid] = arr[i];
    }
#pragma omp parallel sections
    {
#pragma omp section
    {
        merge_sort(left, mid);
    }
#pragma omp section
    {
        merge_sort(right, size - mid);
    }
    }
    merge(arr, left, mid, right, size - mid);
}

int main()
{
    int arr[ARRAY_SIZE];
    int num_threads_array[] = {16};
    int num_threads_array_size = sizeof(num_threads_array) / sizeof(int);

    // Initialize the array with random values
    for (int i = 0; i < ARRAY_SIZE; i++)
    {
        arr[i] = rand() % ARRAY_SIZE;
    }

    // Sort the array using normal merge sort
    clock_t start_time = clock();
    merge_sort(arr, ARRAY_SIZE);
    clock_t end_time = clock();
    double normal_time = ((double)(end_time - start_time)) / CLOCKS_PER_SEC;

    // Sort the array in parallel using OpenMP
    for (int i = 0; i < num_threads_array_size; i++)
    {
        int num_threads = num_threads_array[i];
        printf("Number of threads: %d\n", num_threads);
        start_time = clock();
        omp_set_num_threads(num_threads);
#pragma omp parallel
        {
#pragma omp single
        {
            merge_sort(arr, ARRAY_SIZE);
        }
    }
}

```

```
end_time = clock();
double parallel_time = ((double)(end_time - start_time)) / CLOCKS_PER_SEC;

// Print the time taken by both merge sorts
printf("Time taken (normal merge sort): %f seconds\n", normal_time);
printf("Time taken (parallel merge sort): %f seconds\n", parallel_time);

float speedUp = normal_time / parallel_time;
float efficiency = speedUp / num_threads;

cout << "Speed Up: " << speedUp << "\n";
cout << "Efficiency: " << efficiency << "\n";
printf("\n");
}

return 0;
}
```

Output

```
Windows PowerShell
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> g++ merge.cpp -o merge -fopenmp
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> ./merge
Number of threads: 16
Time taken (normal merge sort): 0.083000 seconds
Time taken (parallel merge sort): 0.070000 seconds
Speed Up: 1.18571
Efficiency: 0.0741071

PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> |
```

Code

```
#include <limits.h>
#include <omp.h>
#include <stdlib.h>
```

```
#include <array>
#include <chrono>
#include <functional>
#include <iostream>
#include <string>
#include <vector>
```

```
using std::chrono::duration_cast;
using std::chrono::high_resolution_clock;
using std::chrono::milliseconds;
using namespace std;
```

```
void s_avg(int arr[], int n)
{
    long sum = 0L;
    int i;
    for (i = 0; i < n; i++)
    {
        sum = sum + arr[i];
    }
    // cout << "\nAverage = " << sum / long(n) << "\n";
}
```

```
void p_avg(int arr[], int n)
{
    long sum = 0L;
    int i;
#pragma omp parallel for reduction(+ : sum) num_threads(16)
    for (i = 0; i < n; i++)
    {
        sum = sum + arr[i];
    }
    // cout << "\nAverage = " << sum / long(n) << "\n";
}
```

```
void s_sum(int arr[], int n)
{
    long sum = 0L;
    int i;
    for (i = 0; i < n; i++)
    {
        sum = sum + arr[i];
    }
    // cout << "\nSum = " << sum << "\n";
}
```

```
void p_sum(int arr[], int n)
```

```

{
    long sum = 0L;
    int i;
#pragma omp parallel for reduction(+ : sum) num_threads(16)
    for (i = 0; i < n; i++)
    {
        sum = sum + arr[i];
    }
    // cout << "\nSum = " << sum << "\n";
}

```

```

void s_max(int arr[], int n)
{
    int max_val = INT_MIN;
    int i;
    for (i = 0; i < n; i++)
    {
        if (arr[i] > max_val)
        {
            max_val = arr[i];
        }
    }
    // cout << "\nMax value = " << max_val << "\n";
}

```

```

void p_max(int arr[], int n)
{
    int max_val = INT_MIN;
    int i;
#pragma omp parallel for reduction(max : max_val) num_threads(16)
    for (i = 0; i < n; i++)
    {
        if (arr[i] > max_val)
        {
            max_val = arr[i];
        }
    }
    // cout << "\nMax value = " << max_val << "\n";
}

```

```

void s_min(int arr[], int n)
{
    int min_val = INT_MAX;
    int i;
    for (i = 0; i < n; i++)
    {
        if (arr[i] < min_val)
        {
            min_val = arr[i];
        }
    }
    // cout << "\nMin value = " << min_val << "\n";
}

```



```

}

void p_min(int arr[], int n)
{
    int min_val = INT_MAX;
    int i;
#pragma omp parallel for reduction(min : min_val) num_threads(16)
    for (i = 0; i < n; i++)
    {
        if (arr[i] < min_val)
        {
            min_val = arr[i];
        }
    }
    // cout << "\nMin value = " << min_val << "\n";
}

int bench_traverse(std::function<void()> traverse_fn)
{
    auto start = high_resolution_clock::now();
    traverse_fn();
    auto stop = high_resolution_clock::now();

    // Subtract stop and start timepoints and cast it to required unit.
    // Predefined units are nanoseconds, microseconds, milliseconds, seconds,
    // minutes, hours. Use duration_cast() function.
    auto duration = duration_cast<milliseconds>(stop - start);

    // To get the value of duration use the count() member function on the
    // duration object
    duration.count();
}

int main(int argc, const char **argv)
{
    if (argc < 2)
    {
        cout << "Specify array length.\n";
        return 1;
    }
    int *a, n, i;

    n = stoi(argv[1]);
    a = new int[n];

    for (int i = 0; i < n; i++)
    {
        a[i] = rand() % n;
    }

    cout << "Generated random array of length " << n << "\n\n";
    omp_set_num_threads(16);

```

```

int sequentialMin = bench_traverse([&
    { s_min(a, n); });

int parallelMin = bench_traverse([&
    { p_min(a, n); });

int sequentialMax = bench_traverse([&
    { s_max(a, n); });

int parallelMax = bench_traverse([&
    { p_max(a, n); });

int sequentialSum = bench_traverse([&
    { s_sum(a, n); });

int parallelSum = bench_traverse([&
    { p_sum(a, n); });

int sequentialAverage = bench_traverse([&
    { s_avg(a, n); });

int parallelAverage = bench_traverse([&
    { p_avg(a, n); });

cout << "Sequential Min: " << sequentialMin << "ms\n";
cout << "Parallel (16) Min: " << parallelMin << "ms\n";
cout << "Speed Up for Min: " << (float)sequentialMin / parallelMin << "\n";
cout << "Efficiency for Min: " << ((float)sequentialMin / parallelMin) / 16 << "\n";

cout << "\nSequential Max: " << sequentialMax << "ms\n";
cout << "Parallel (16) Max: " << parallelMax << "ms\n";
cout << "Speed Up for Max: " << (float)sequentialMax / parallelMax << "\n";
cout << "Efficiency for Max: " << ((float)sequentialMax / parallelMax) / 16 << "\n";

cout << "\nSequential Sum: " << sequentialSum << "ms\n";
cout << "Parallel (16) Sum: " << parallelSum << "ms\n";
cout << "Speed Up for Sum: " << (float)sequentialSum / parallelSum << "\n";
cout << "Efficiency for Sum: " << ((float)sequentialSum / parallelSum) / 16 << "\n";

cout << "\nSequential Average: " << sequentialAverage << "ms\n";
cout << "Parallel (16) Average: " << parallelAverage << "ms\n";
cout << "Speed Up for Average: " << (float)sequentialAverage / parallelAverage << "\n";
cout << "Efficiency for Average: " << ((float)sequentialAverage / parallelAverage) / 16 << "\n";

return 0;
}

```

Output

```
Windows PowerShell
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> g++ reduction.cpp -o reduction -fopenmp
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> ./reduction 10000000
Generated random array of length 10000000

Sequential Min: 36ms
Parallel (16) Min: 7ms
Speed Up for Min: 5.14286
Efficiency for Min: 0.321429

Sequential Max: 36ms
Parallel (16) Max: 4ms
Speed Up for Max: 9
Efficiency for Max: 0.5625

Sequential Sum: 36ms
Parallel (16) Sum: 4ms
Speed Up for Sum: 9
Efficiency for Sum: 0.5625

Sequential Average: 37ms
Parallel (16) Average: 5ms
Speed Up for Average: 7.4
Efficiency for Average: 0.4625
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\1> |
```

Addition of two Large Vectors

Code

```
#include <stdio.h>
```

```
__global__ void vectorAdd(float *a, float *b, float *c, int n)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < n)
    {
        c[i] = a[i] + b[i];
    }
}

int main()
{
    int n = 1000000;
    size_t bytes = n * sizeof(float);

    // Allocate memory on the host
    float *h_a = (float *)malloc(bytes);
    float *h_b = (float *)malloc(bytes);
    float *h_c = (float *)malloc(bytes);

    // Initialize the vectors
    for (int i = 0; i < n; i++)
    {
        h_a[i] = i;
        h_b[i] = i + 1;
    }

    // Allocate memory on the device
    float *d_a, *d_b, *d_c;
    cudaMalloc(&d_a, bytes);
    cudaMalloc(&d_b, bytes);
    cudaMalloc(&d_c, bytes);

    // Copy data from host to device
    cudaMemcpy(d_a, h_a, bytes, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, h_b, bytes, cudaMemcpyHostToDevice);

    // Launch kernel on the device
    int threadsPerBlock = 256;
    int blocksPerGrid = (n + threadsPerBlock - 1) / threadsPerBlock;
    vectorAdd<<<blocksPerGrid, threadsPerBlock>>>>(d_a, d_b, d_c, n);

    // Copy result from device to host
    cudaMemcpy(h_c, d_c, bytes, cudaMemcpyDeviceToHost);

    // Print first 10 elements of both vectors
    printf("First 10 elements of vector a:\n");
    for (int i = 0; i < 10; i++)
    {
```

```

    printf("%.2f ", h_a[i]);
}
printf("\n");
printf("Size of vector a: %d\n", n);
printf("\n");

printf("First 10 elements of vector b:\n");
for (int i = 0; i < 10; i++)
{
    printf("%.2f ", h_b[i]);
}
printf("\n");
printf("Size of vector b: %d\n", n);
printf("\n");

// Print first 10 elements of resultant vector
printf("First 10 elements of resultant vector:\n");
for (int i = 0; i < 10; i++)
{
    printf("%.2f ", h_c[i]);
}
printf("\n");

// Print size of resultant vector
printf("Size of resultant vector: %d\n", n);

// Free memory
free(h_a);
free(h_b);
free(h_c);
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);

return 0;
}

```

Output

```
Windows PowerShell
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\4> nvcc addition.cu -o addition
addition.cu
tmpxft_000035f4_00000000-10_addition.cudafe1.cpp
  Creating library addition.lib and object addition.exp
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\4> ./addition
First 10 elements of vector a:
0.00 1.00 2.00 3.00 4.00 5.00 6.00 7.00 8.00 9.00
Size of vector a: 1000000

First 10 elements of vector b:
1.00 2.00 3.00 4.00 5.00 6.00 7.00 8.00 9.00 10.00
Size of vector b: 1000000

First 10 elements of resultant vector:
1.00 3.00 5.00 7.00 9.00 11.00 13.00 15.00 17.00 19.00
Size of resultant vector: 1000000
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\4> |
```

Matrix Multiplication using CUDA C Code

```
#include <stdio.h>
```

```
#define TILE_WIDTH 32
```

```
__global__ void matrixMul(float *a, float *b, float *c, int m, int n, int p)
```

```
{  
    __shared__ float As[TILE_WIDTH][TILE_WIDTH];  
    __shared__ float Bs[TILE_WIDTH][TILE_WIDTH];
```

```
  
    int bx = blockIdx.x;  
    int by = blockIdx.y;  
    int tx = threadIdx.x;  
    int ty = threadIdx.y;
```

```
  
    int row = by * TILE_WIDTH + ty;  
    int col = bx * TILE_WIDTH + tx;
```

```
  
    float Cvalue = 0.0;
```

```
  
    for (int k = 0; k < n / TILE_WIDTH; k++)  
    {  
        As[ty][tx] = a[row * n + k * TILE_WIDTH + tx];  
        Bs[ty][tx] = b[(k * TILE_WIDTH + ty) * p + col];
```

```
  
        __syncthreads();
```

```
  
        for (int i = 0; i < TILE_WIDTH; i++)  
        {  
            Cvalue += As[ty][i] * Bs[i][tx];  
        }
```

```
  
        __syncthreads();  
    }
```

```
  
    c[row * p + col] = Cvalue;  
}
```

```
  
int main()  
{  
    int m = 1024;  
    int n = 1024;  
    int p = 1024;  
    size_t bytesA = m * n * sizeof(float);  
    size_t bytesB = n * p * sizeof(float);  
    size_t bytesC = m * p * sizeof(float);
```

```
  
    // Allocate memory on the host  
    float *h_a = (float *)malloc(bytesA);  
    float *h_b = (float *)malloc(bytesB);  
    float *h_c = (float *)malloc(bytesC);
```

```

// Initialize matrices
for (int i = 0; i < m * n; i++)
{
    h_a[i] = 1.0;
}
for (int i = 0; i < n * p; i++)
{
    h_b[i] = 2.0;
}

// Allocate memory on the device
float *d_a, *d_b, *d_c;
cudaMalloc(&d_a, bytesA);
cudaMalloc(&d_b, bytesB);
cudaMalloc(&d_c, bytesC);

// Copy data from host to device
cudaMemcpy(d_a, h_a, bytesA, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, h_b, bytesB, cudaMemcpyHostToDevice);

// Launch kernel on the device
dim3 dimBlock(TILE_WIDTH, TILE_WIDTH);
dim3 dimGrid((p + dimBlock.x - 1) / dimBlock.x, (m + dimBlock.y - 1) / dimBlock.y);
matrixMul<<<dimGrid, dimBlock>>>(d_a, d_b, d_c, m, n, p);

// Copy result from device to host
cudaMemcpy(h_c, d_c, bytesC, cudaMemcpyDeviceToHost);

// Print 3x3 parts of both matrices
printf("Matrix A (3x3 part):\n");
for (int i = 0; i < 3; i++)
{
    for (int j = 0; j < 3; j++)
    {
        printf("%.2f ", h_a[i * n + j]);
    }
    printf("\n");
}
printf("Size of Matrix A: %dx%d\n", m, n);
printf("\n");

printf("Matrix B (3x3 part):\n");
for (int i = 0; i < 3; i++)
{
    for (int j = 0; j < 3; j++)
    {
        printf("%.2f ", h_b[i * p + j]);
    }
    printf("\n");
}
printf("Size of Matrix B: %dx%d\n", n, p);

```



```

printf("\n");

// Print 3x3 part of resultant matrix
printf("Resultant Matrix (3x3 part):\n");
for (int i = 0; i < 3; i++)
{
    for (int j = 0; j < 3; j++)
    {
        printf("%.2f ", h_c[i * p + j]);
    }
    printf("\n");
}

// Print size of resultant matrix
printf("Size of Resultant Matrix: %dx%d\n", m, p);

// Free memory on the host and device
free(h_a);
free(h_b);
free(h_c);
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);

return 0;
}

```

Output

```
Windows PowerShell
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\4> nvcc multiplication.cu -o multiplication
multiplication.cu
tmpxft_00006360_00000000-10_multiplication.cudafe1.cpp
Creating library multiplication.lib and object multiplication.exp
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\4> ./multiplication
Matrix A (3x3 part):
1.00 1.00 1.00
1.00 1.00 1.00
1.00 1.00 1.00
Size of Matrix A: 1024x1024

Matrix B (3x3 part):
2.00 2.00 2.00
2.00 2.00 2.00
2.00 2.00 2.00
Size of Matrix B: 1024x1024

Resultant Matrix (3x3 part):
2048.00 2048.00 2048.00
2048.00 2048.00 2048.00
2048.00 2048.00 2048.00
Size of Resultant Matrix: 1024x1024
PS D:\Gaurav\College\4 BE\Sem 8\High Performance Computing\Lab Assignments\4> |
```